

## Week 5

### Deep Learning II

Convolutional Neural Networks

Dr Anagi Gamachchi

Discipline of Information Systems and Business Analytics,  
Deakin Business School



# AI Application at Woolworth

[NEWS](#)[ANALYSIS](#)[MARKETS](#)[RETAIL](#)[SUPPLY](#)[TECHNOLOGY](#)[PACKAGING](#)[EVENTS](#)[SUBSCRIBE](#)[ADVERT](#)[PRODUCE PLUS](#)

## Woolworths scans fresh produce with AI

By Liam O'Callaghan | 18 February 2020



Woolworths is aiming to speed up scanning for shoppers with a new artificial intelligence (AI) powered scale that automatically identifies loose fresh produce.



The scales, developed by Australian company Tiliter, are currently being trialled in three Sydney stores in conjunction with the retailer's Scan&Go app.

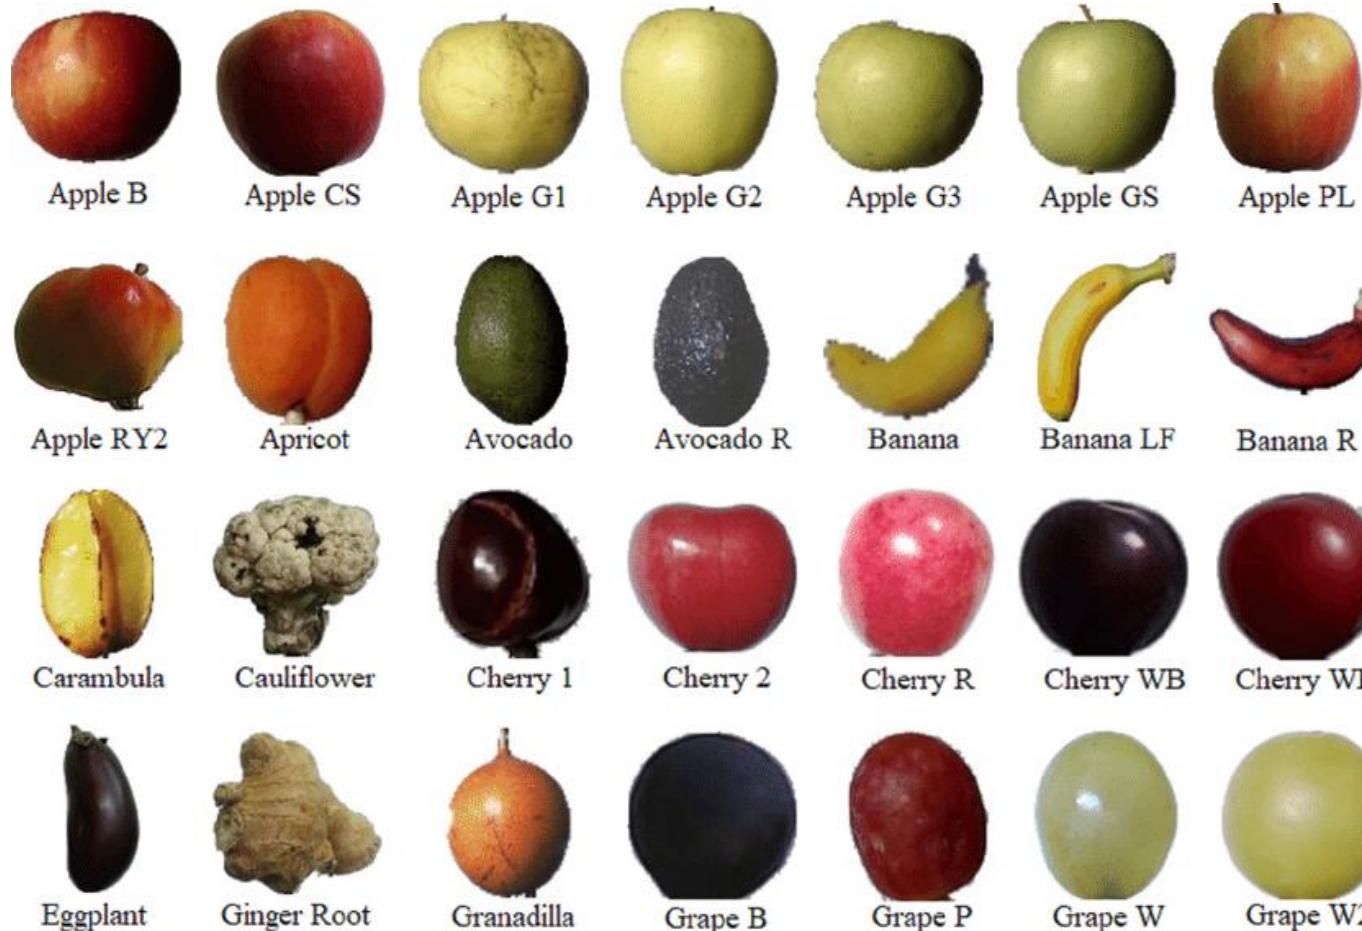
Consumers can place their fruit or vegetables on the scale surface and then the technology will detect the type of produce and identify the variety.



Woolworths is aiming to speed up scanning for shoppers with a new artificial intelligence (AI) powered scale that automatically identifies loose fresh produce.

# Discussion Question

How can Machine identify the correct type of fruits?



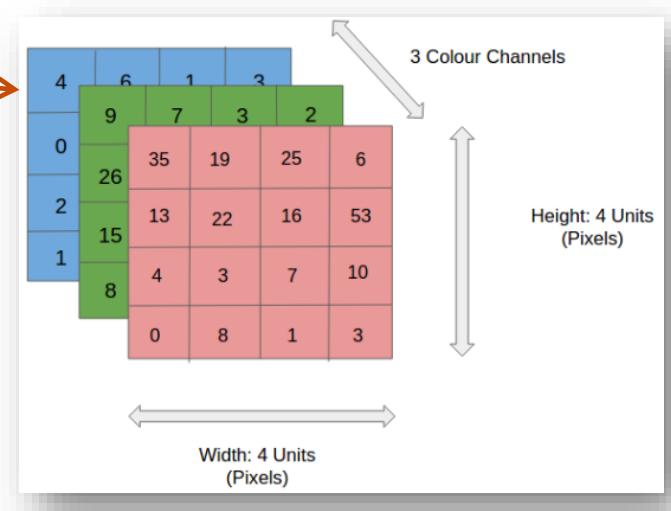
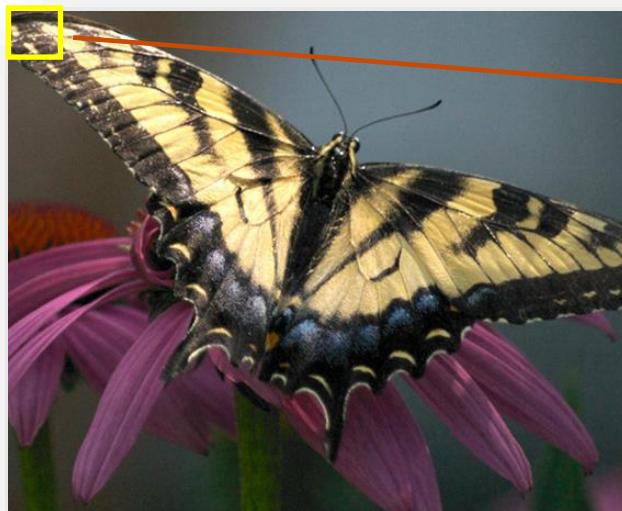
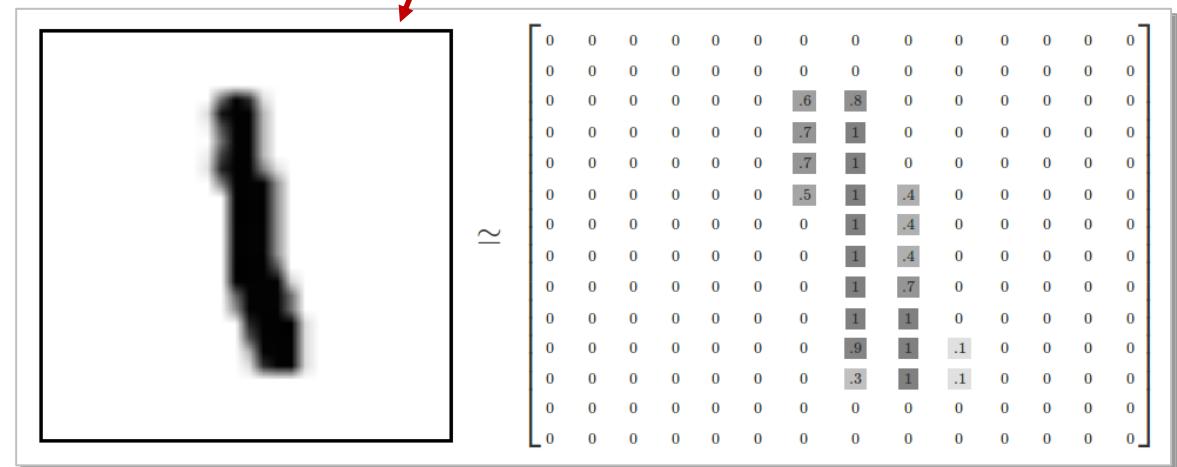
# Image Representation



**Representation of visual data (images) is based on matrices of pixels**

**Pixel:** Is the smallest unit of an image

- Has an address in the matrix (row, column)
  - Has an intensity...
    - **Greyscale images:** 0 for black and 255 for white
    - **RGB images:** 256 shades of red, green, and blue colors



**Resolution of the screen** (image container): Is the size of the matrix in terms of the number of rows and columns (pixels)

# How do Convolutional Neural Networks (CNN) work?

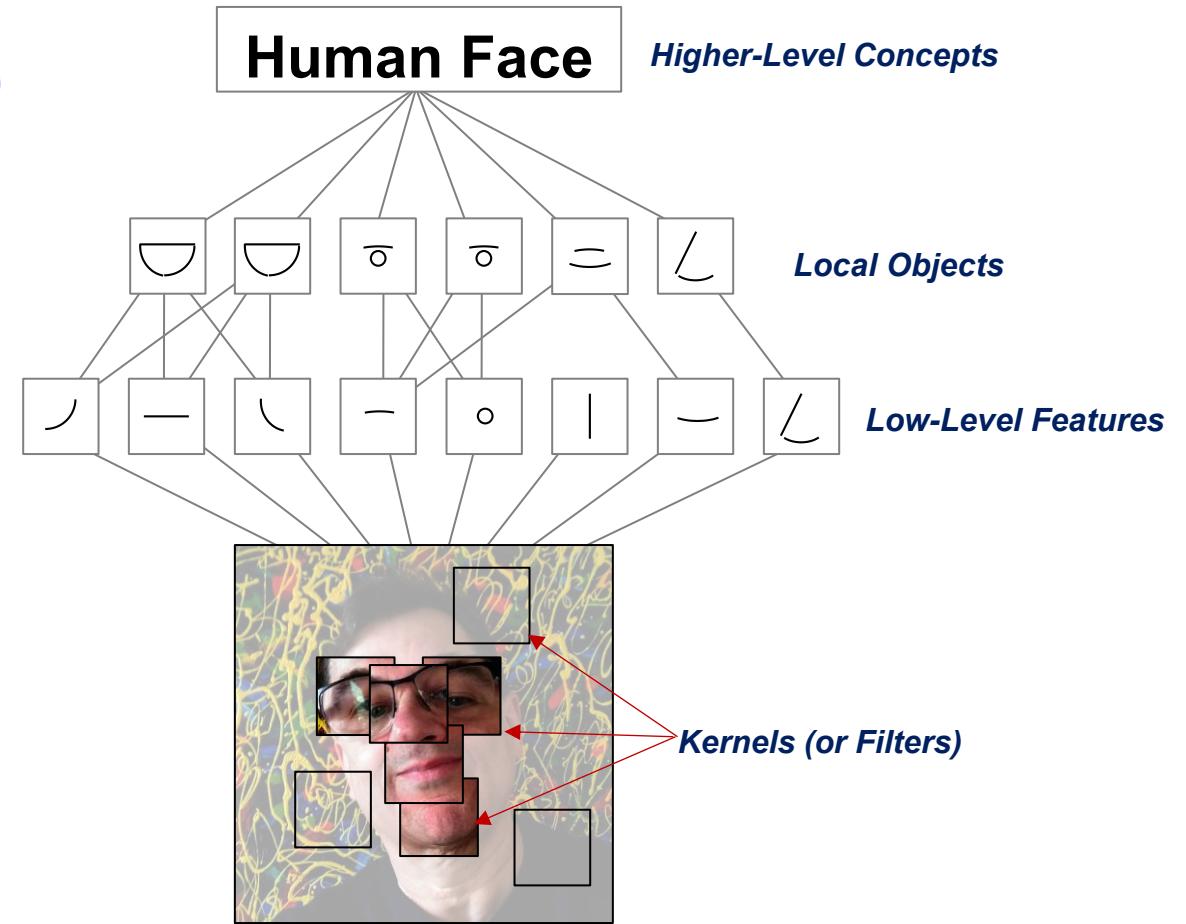
The Deep Learning magic is in meaning (**representation**) forming at different neural network layers

Images are scanned by tiny **filters** (also called **kernels**).

They are responsible for matching and identifying **low-level visual features**, such as edges of different direction, shades or colours.

Once detected at the next level these features are being combined into **local objects**, such as eyes, mouths, noses or glasses.

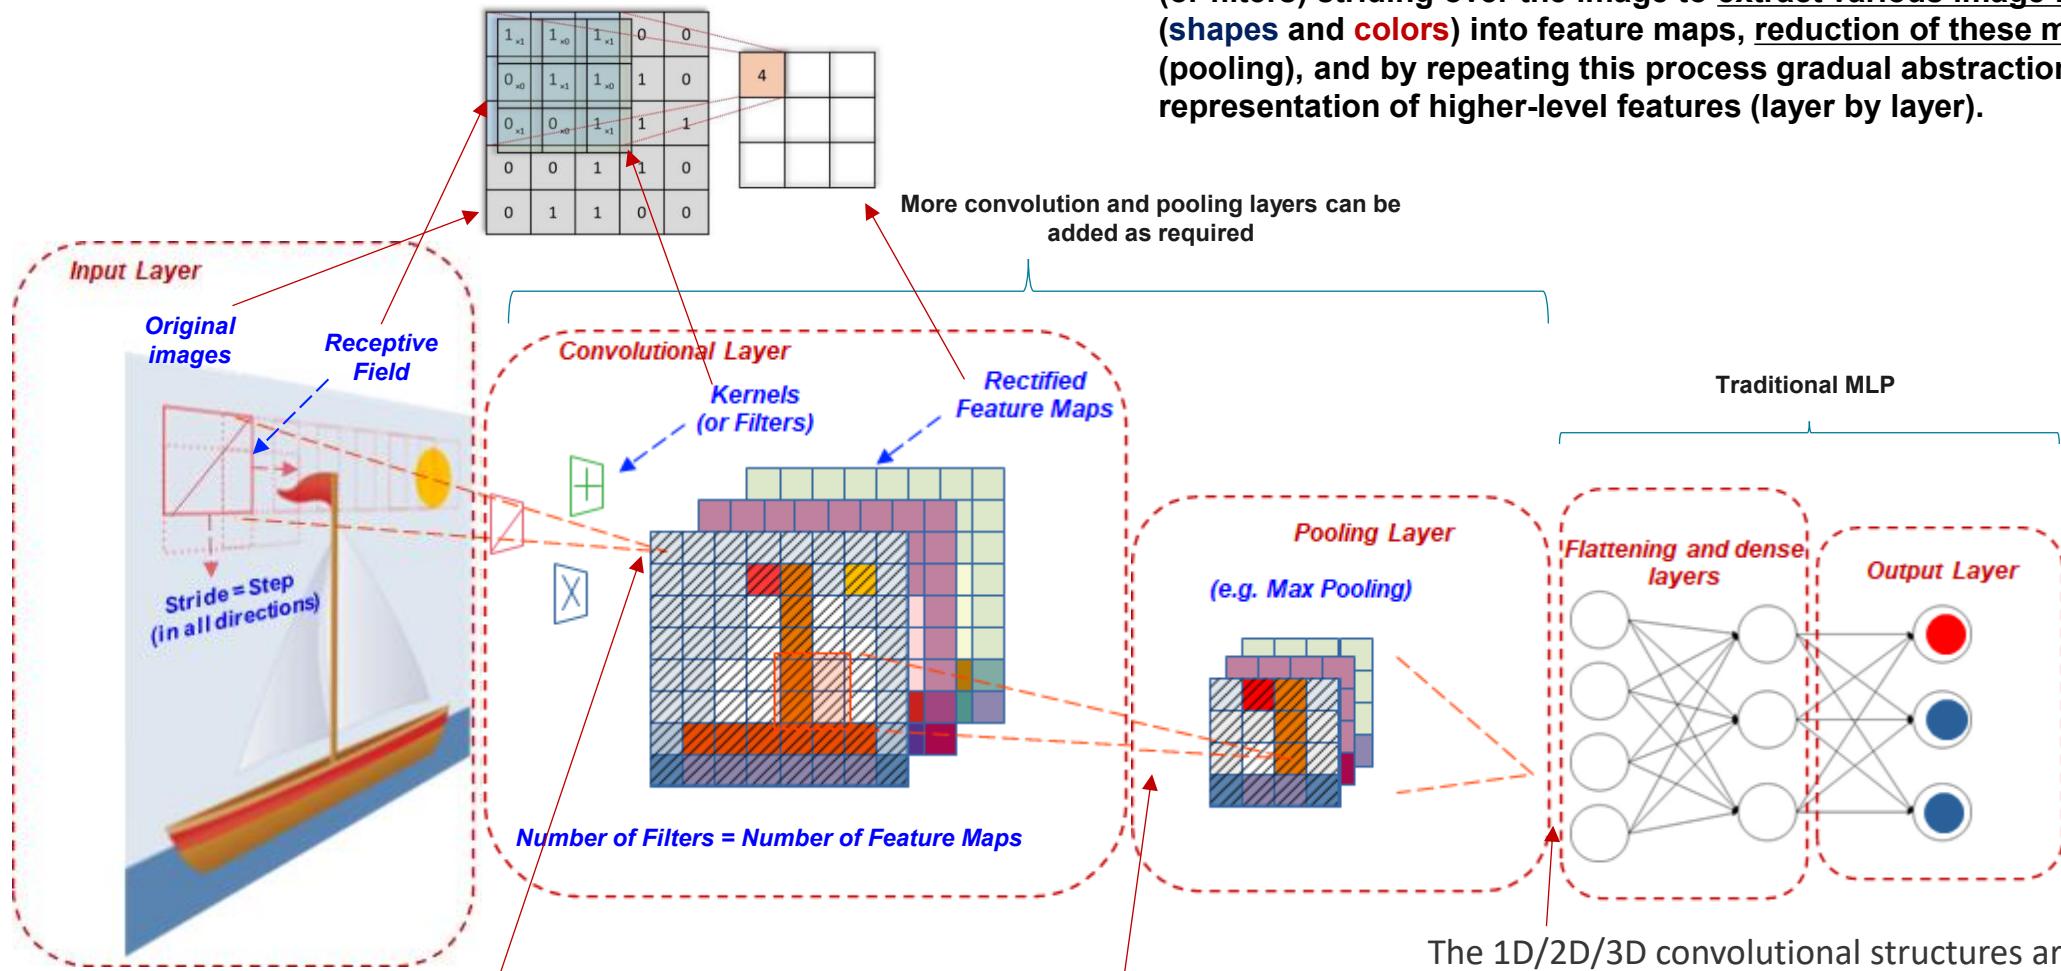
These local objects in turn combine into **high-level concepts**, such as a “human face”.



A **kernel** is a numeric matrix, used to identify lines and edges in different direction, sharpening or blurring lines, identifying colours, etc. Kernels are discovered automatically.



# Architecture of CNN

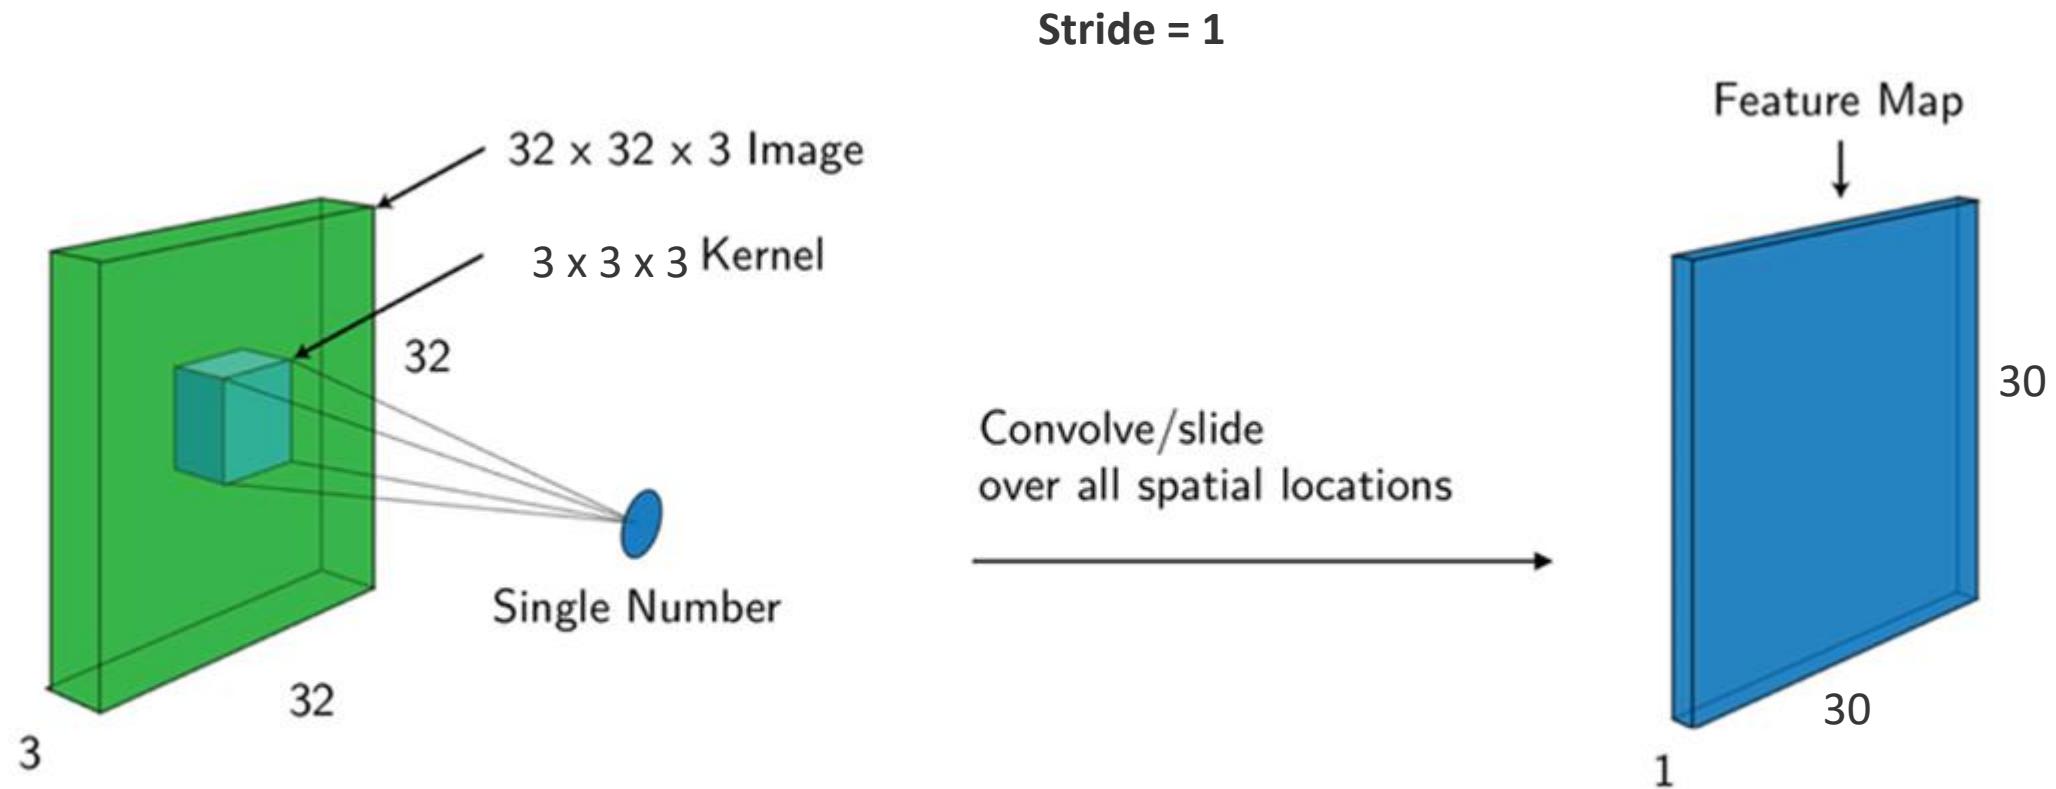


The **convolution operation** is the weighted sum of the receptive field and the kernel (usually 3x3 or 5x5), across all channels (RGB).

**Max-pooling operation** takes the maximum value from each channel within a window (usually 2x2) of the feature map, which is sliding over it (with a stride/step of 2 pixels), thus reducing the image dimension (by half).

**CNNs** link different types of layers which aim at using small kernels (or filters) striding over the image to extract various image features (shapes and colors) into feature maps, reduction of these maps (pooling), and by repeating this process gradual abstraction and representation of higher-level features (layer by layer).

# Convolution Layer



# Max Pooling

Rectified Feature Map

10	20	25	70
8	11	80	40
4	5	8	9
6	7	10	11

Max pooling with 2\*2 filters and stride 2

$$\text{Max}(4,5,6,7)=7$$

Pooled Feature map

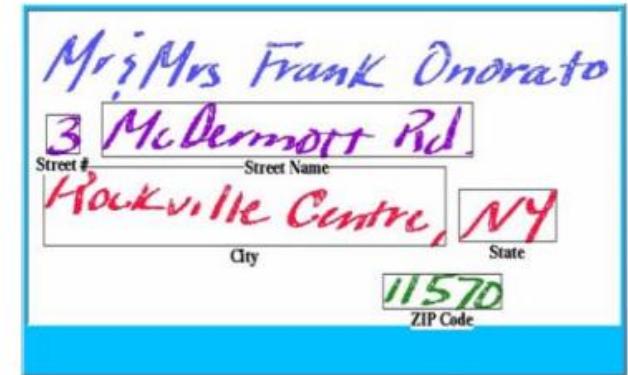
20	80
7	11

# CNN Applications

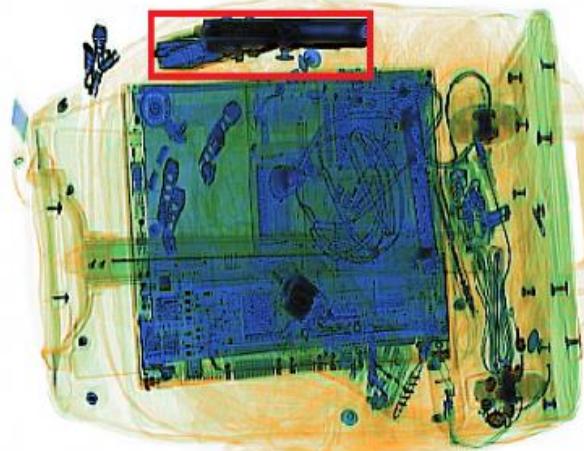
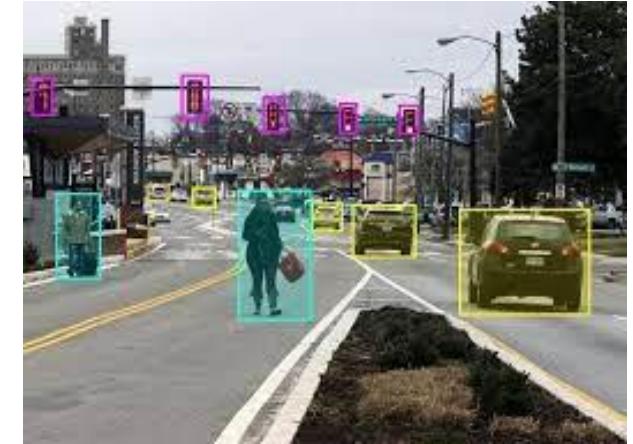
## Applications of CNN include:

- Object Detection (e.g., fruits)
- Hand-writing recognition (e.g. transcription)
- Facial recognition (e.g. social media tagging)
- Emotion recognition (e.g. in lie detectors)
- Scene labelling (e.g. self-driving cars)
- Action recognition (e.g. suspicious behaviour)
- Motion detection (e.g. security cameras)
- MRI / CT diagnosis (e.g. cardiac imaging)
- Satellite image processing (e.g. for planning)
- Land feature recognition (e.g. missile systems)
- Echo-sound processing (e.g. oil prospecting)
- Object detection (e.g. airport scanners)
- Object tracking (e.g. players/ball in sport)
- Colouring and noise removal (e.g. photography)
- ...

Handwriting  
recognition for  
postal automation



Scene labelling  
self-driving car



Threaten Object-Detection  
in Bag Scanning at Airport



# CNN in Python

## ❑ Data Set:

**CIFAR10 data sets** includes images of consists of 60000 32x32 colour images in 10 classes, with 6000 images per class

## Your task is to:

*Recognize the corresponding object in the images.*

## ❑ Tool:

Python + Tensorflow (with Keras)

## ❑ Method:

*Convolutional Neural Nets (CNN)*

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



# CNN in Python

## Data Loading

```
from tensorflow.keras.datasets import cifar10

# Data parameters
img_rows, img_cols = 32, 32
channels = 3

num_classes = 10
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

plot_images(x_train[0:20], cols=5, figsize=[7,7])
```



11

## Data Preparation

```
x_train /= 255  
x_test /= 255
```

Why?

```
# convert class vectors to binary class matrices
y_train = to_categorical(y_train, num_classes)
y_test = to_categorical(y_test, num_classes)

print('Train shape: x=', x_train.shape, ', y=', y_train.shape)
print('Test shape: x=', x_test.shape, ', y=', y_test.shape)
```

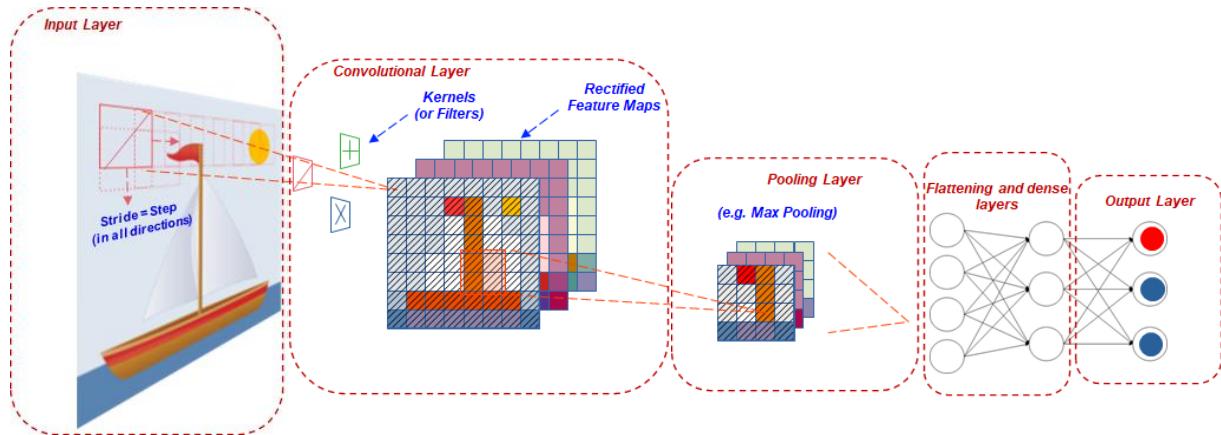
Train shape: x= (50000, 32, 32, 3) , y= (50000, 10)  
Test shape: x= (10000, 32, 32, 3) , y= (10000, 10)

Why 3?

Why 10?

# CNN in Python (cont.)

*Deep Learning CNN model is created*



```
def model_2():
    model = Sequential()
    model.add(Conv2D(32, kernel_size=(3, 3),
                    strides=(1, 1),
                    activation='relu',
                    input_shape=(img_rows, img_cols, channels)))
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    Reduced to a quarter
    model.add(Dropout(0.25))  
    Prevent overfitting,
    model.add(Flatten())  
    randomly dropping
    model.add(Dense(128, activation='relu'))  
    some connections with
    model.add(Dropout(0.5))  
    preceding layer
    model.add(Dense(num_classes, activation='softmax'))
    model.summary()
    return model
```

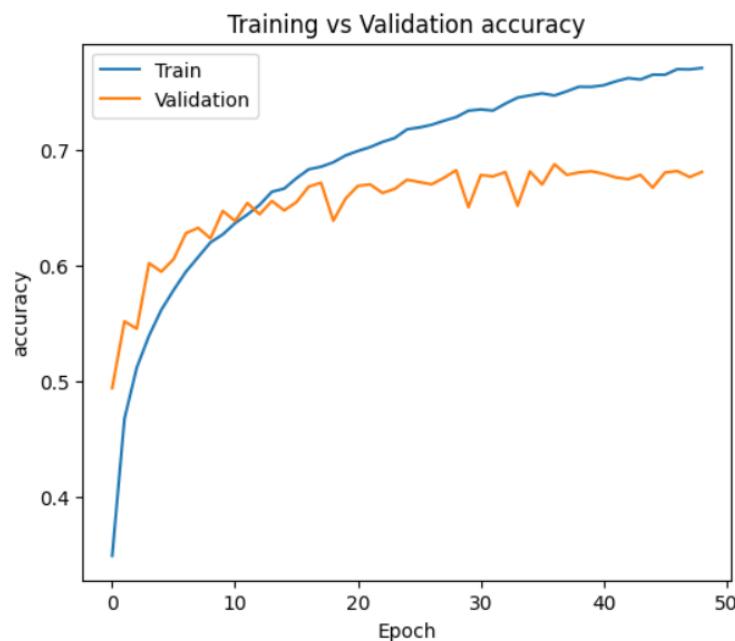
Annotations for the code:

- 1 Convolution and 1 Pooling Layers: Brackets group the first two layers of the model.
- 1 Flattening and 1 Dense Layers: Brackets group the third and fourth layers of the model.
- 1 Output Layer: Brackets group the final layer of the model.

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 30, 30, 32)	896 = $(3 \times 3 \times 3 + 1) * 32$
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_2 (Dropout)	(None, 15, 15, 32)	0
flatten_1 (Flatten)	(None, 7200) $= 15 \times 15 \times 32$	0
dense_2 (Dense)	(None, 128)	921728 = $(7200 + 1) * 128$
dropout_3 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290 = $(128 + 1) * 10$
<hr/>		
Total params: 923,914		
Trainable params: 923,914		
Non-trainable params: 0		

# CNN in Python (cont.)

- We will experiment with various *optimizers* or algorithms searching for the best set of network *weights* and *biases*
- The error function is called *loss* to guide the optimiser
- Other *metrics* can also be used measure the net performance, e.g. accuracy



The training process terminated after 49 epochs.

13

```
model.compile(loss=categorical_crossentropy,  
              optimizer=RMSprop(learning_rate=0.001, weight_decay=1e-6),  
              metrics='accuracy') ← optimizers  
  
hist = model.fit(x_train, y_train,  
                  batch_size=128,  
                  epochs=100, ← Train the model  
                  verbose=2,  
                  validation_data=(x_test, y_test),  
                  validation_split=0.2,  
                  callbacks=keras_callbacks)
```

Epoch 1/100

391/391 - 3s - loss: 2.1041 - accuracy: 0.2456 - val\_loss: 1.8539 - val\_accuracy: 0.3309

Epoch 2/100

391/391 - 2s - loss: 1.8945 - accuracy: 0.3132 - val\_loss: 1.7669 - val\_accuracy: 0.3608

Epoch 3/100

Adadelta(lr=0.001, rho=0.95, epsilon=1e-07)  
Adadelta(lr=0.05, rho=0.99, epsilon=1e-07)  
Adam(lr=0.001, beta\_1=0.9, beta\_2=0.999, epsilon=1e-07)

Other Optimizers

Final Performance measures:

Train loss: 0.2722  
Train accuracy: 0.9236

Test loss: 1.098  
Test accuracy: 0.6806



# CNN in Python (cont.)

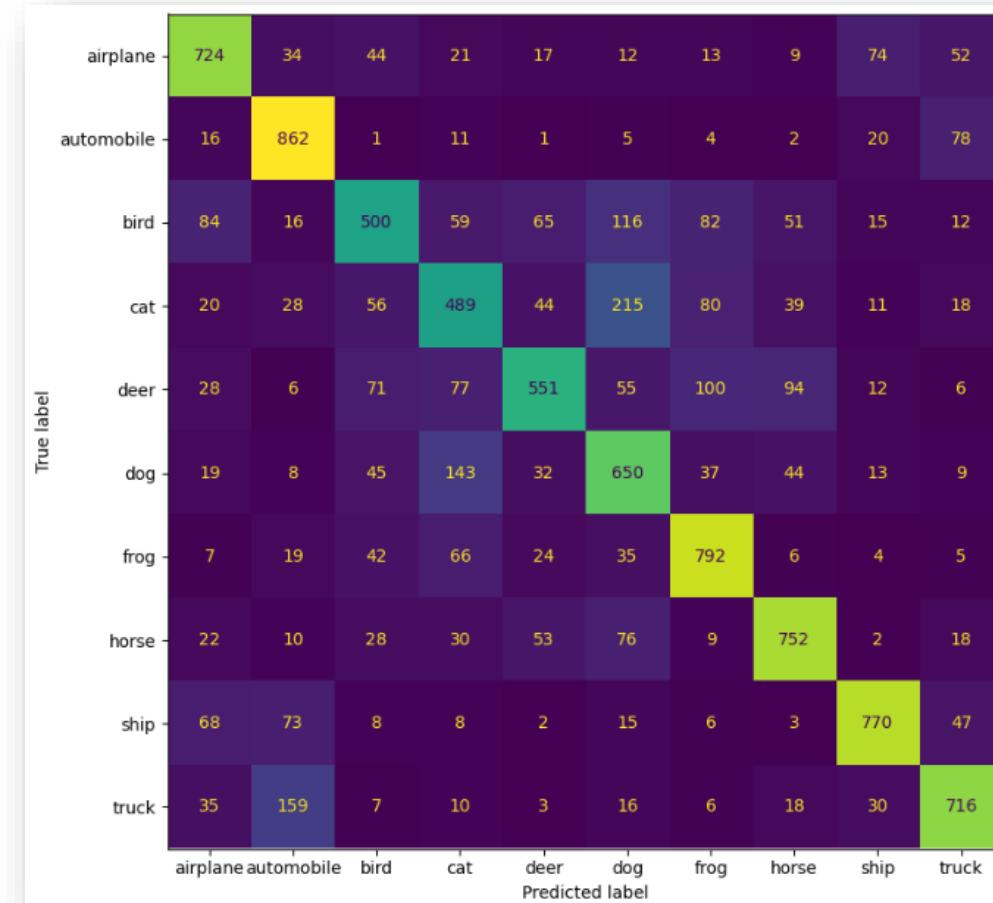
## Evaluation Metrics

The result of Kappa is : 0.645

The result of the classification report is:

	precision	recall	f1-score	support
airplane	0.71	0.72	0.72	1000
automobile	0.71	0.86	0.78	1000
bird	0.62	0.50	0.55	1000
cat	0.54	0.49	0.51	1000
deer	0.70	0.55	0.61	1000
dog	0.54	0.65	0.59	1000
frog	0.70	0.79	0.74	1000
horse	0.74	0.75	0.75	1000
ship	0.81	0.77	0.79	1000
truck	0.75	0.72	0.73	1000
accuracy			0.68	10000
macro avg	0.68	0.68	0.68	10000
weighted avg	0.68	0.68	0.68	10000

Confusion Matrix



# Discussion

```
def model_1():
    model = Sequential()
    model.add(Flatten(input_shape=(img_rows, img_cols, channels)))
    model.add(Dense(128, activation='relu'))
    model.add(Dropout(0.2))
    model.add(Dense(10, activation='softmax'))
    model.summary()
    return model
```

Is this a CNN model?

The result of Kappa is : 0.387				
The result of the classification report is:				
	precision	recall	f1-score	support
airplane	0.60	0.36	0.46	1000
automobile	0.55	0.58	0.56	1000
bird	0.37	0.29	0.33	1000
cat	0.30	0.28	0.29	1000
deer	0.47	0.29	0.36	1000
dog	0.42	0.30	0.35	1000
frog	0.40	0.66	0.50	1000
horse	0.53	0.48	0.51	1000
ship	0.59	0.56	0.58	1000
truck	0.38	0.66	0.48	1000
accuracy			0.45	10000
macro avg	0.46	0.45	0.44	10000
weighted avg	0.46	0.45	0.44	10000

```
def model_2():
    model = Sequential()
    model.add(Conv2D(32, kernel_size=(3, 3),
                    activation='relu',
                    input_shape=(img_rows, img_cols, channels)))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.25))
    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(num_classes, activation='softmax'))
    model.summary()
    return model
```

Which model would perform better? And Why?

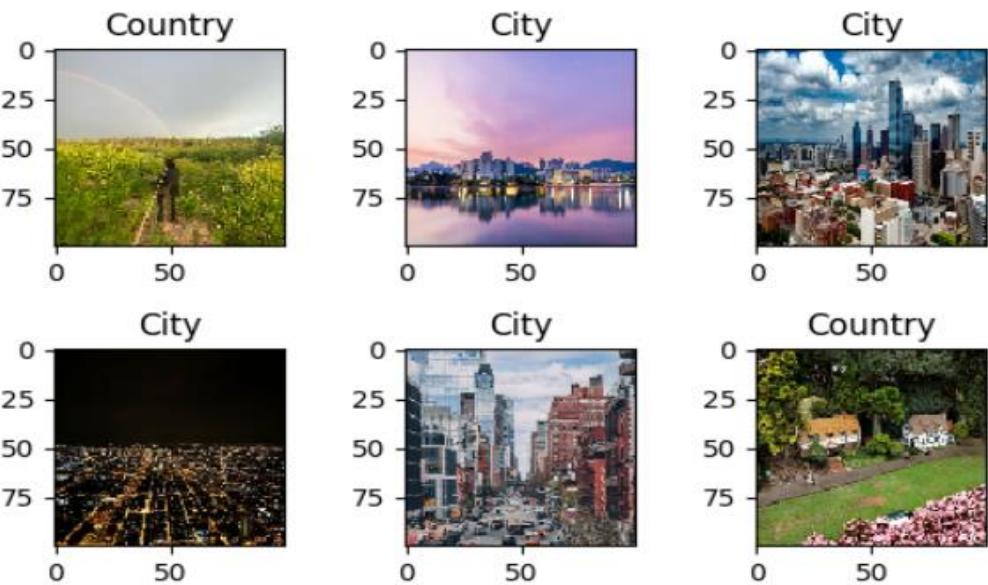
The result of Kappa is : 0.645				
The result of the classification report is:				
	precision	recall	f1-score	support
airplane	0.71	0.72	0.72	1000
automobile	0.71	0.86	0.78	1000
bird	0.62	0.50	0.55	1000
cat	0.54	0.49	0.51	1000
deer	0.70	0.55	0.61	1000
dog	0.54	0.65	0.59	1000
frog	0.70	0.79	0.74	1000
horse	0.74	0.75	0.75	1000
ship	0.81	0.77	0.79	1000
truck	0.75	0.72	0.73	1000
accuracy			0.68	10000
macro avg	0.68	0.68	0.68	10000
weighted avg	0.68	0.68	0.68	10000

How to improve prediction performance?

# Working with Real Digital Photos

```
drive.mount('/content/drive')

# Set the paths to the folders containing the image files
city_path = '/content/drive/MyDrive/Colab Notebooks/dataset/city'
country_path = '/content/drive/MyDrive/Colab Notebooks/dataset/country'
```



```
# Iterate through the files in the first folder
for file in os.listdir(city_path):
    # Check if the file is a jpeg or jpg file
    if file.endswith('.jpeg') or file.endswith('.jpg'):
        # Load the image data from the file using TensorFlow
        img = tf.io.read_file(os.path.join(city_path, file))
        img = tf.image.decode_jpeg(img)
        img = tf.image.resize(img, (100, 100))
        # Assign a label to the file
        label = 'City'
        # Add the image data and label to the data list
        data.append((img, label))
```

## In this lecture, we have covered:

- The concepts and architectures of CNN.
- Experiments with various model architectures of CNN in Python.
- Discussion on the applications of these deep learning techniques.

## Summary