

Supervised Learning

Week 8

KNN

Decision Tree (DT)

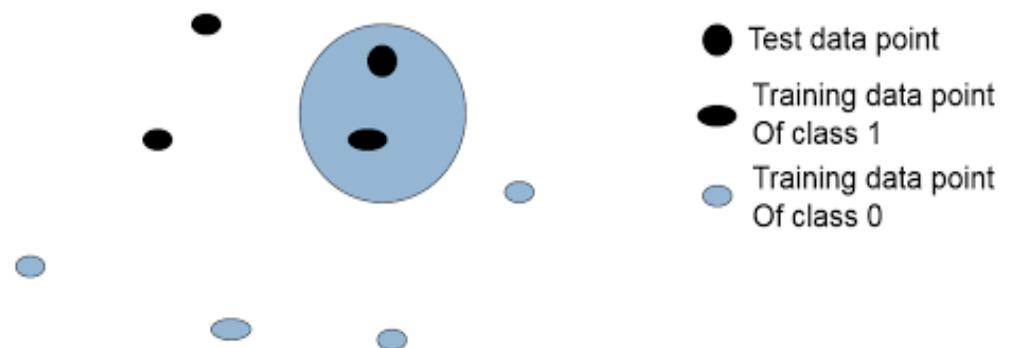
KNN

Algorithm & Variants

Best K

KNN Algorithm & Variants

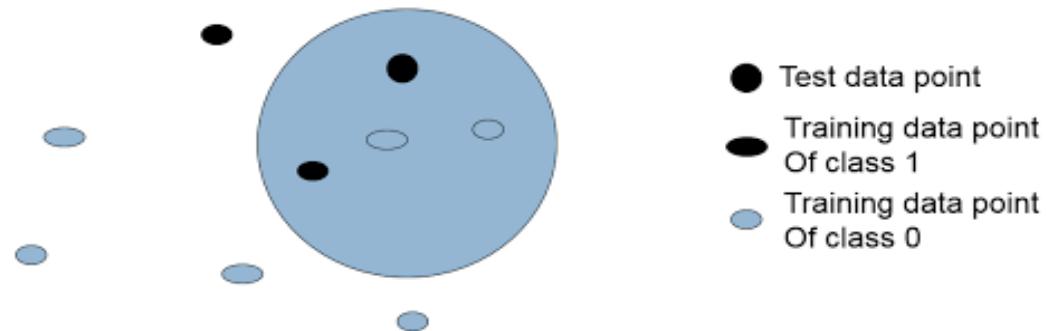
- For both classification and regression
 - A useful technique is to **assign weights based on the neighbours.**
 - The **nearer neighbors contribute more** to the average than the **distant ones.**
- The basic idea is to label the test data point same as the nearest neighbor.



KNN

Algorithm & Variants...

- How many neighbors?
 - K=?



- Let's say someone would like to **check K nearest neighbours** of the test point to make the decision.
 - **Label a test instance** same as the **majority label of the K-nearest** neighbours.
- The figure is an example of **3-NN** classification.

KNN

Algorithm & Variants...

- How to make the majority decisions?
 - Mode of the class labels
 - Discrete cases
 - Average or mean distances
 - Continuous cases
 - Distance-weighted nearest neighbour algorithm (Shepard's method)
 - Assign weights to the neighbours based on their distance from the test point.
 - Weight may be inverse square of the distances ($1/D^2$)
 - Higher the distance of the neighbour, lower its weight.

KNN

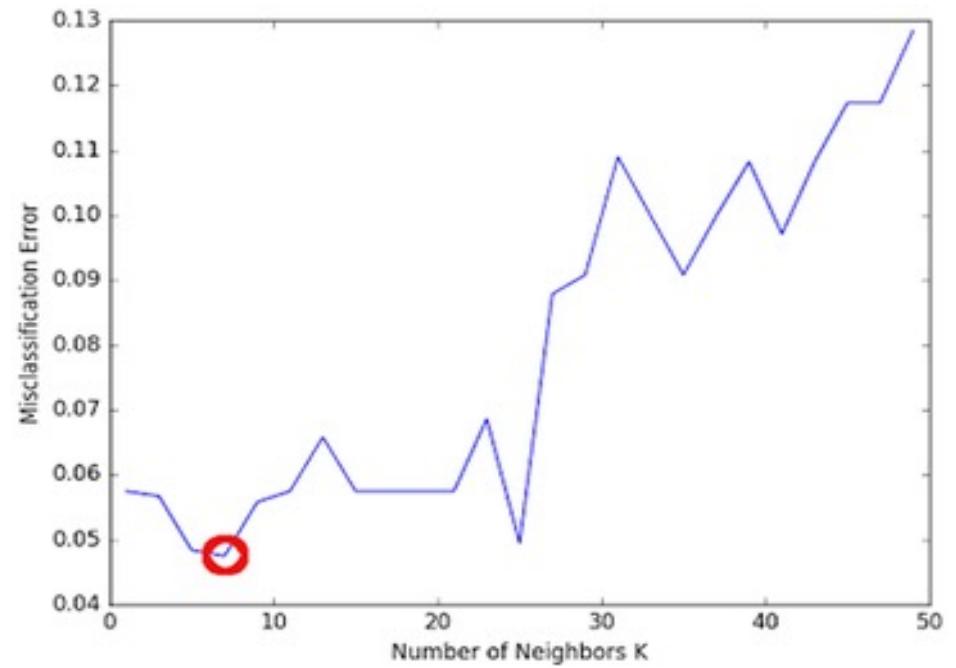
Best number of neighbors (K)

- What is the **importance of the variable K?**
 - K controls the **shape of the decision boundary**
- Small values of K
 - Restrains the region of a given prediction
 - Forces classifier to be more **focused on the close regions and neighbours**
 - This will result in a **low bias** and **high variance**
- Higher values of K
 - Asking for more **information from distant training points**
 - Smoother decision boundaries
 - **Lower variance but increases bias**

KNN

Best number of neighbors (K)...

- Finding the best K
 - There is **no rule of thumb** in selecting K_{max} since it depends on your desired rate of exploration for K
 - A simple and handy method
 - **Cross-validation** to partition your data into test and training samples
 - **Evaluate model** with different ranges of K values
 - $$K = 1, \dots, K_{max}$$
 - The misclassification error can be used as a measurement of performance



Remarks

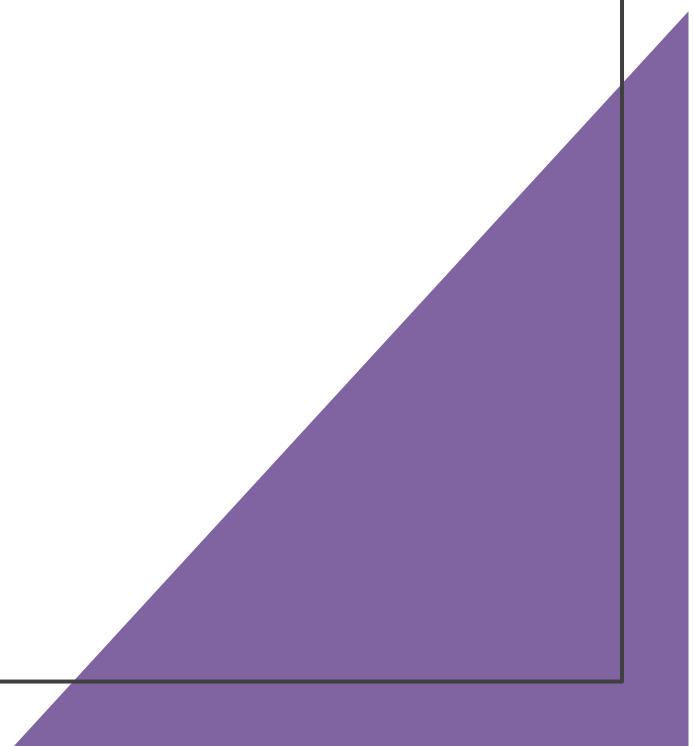
- Learning is very **simple** (actually, no learning involved).
- **Classification is very time consuming** because we need to find distance with all the training instances.

Decision Tree (DT)

Regression/classification trees

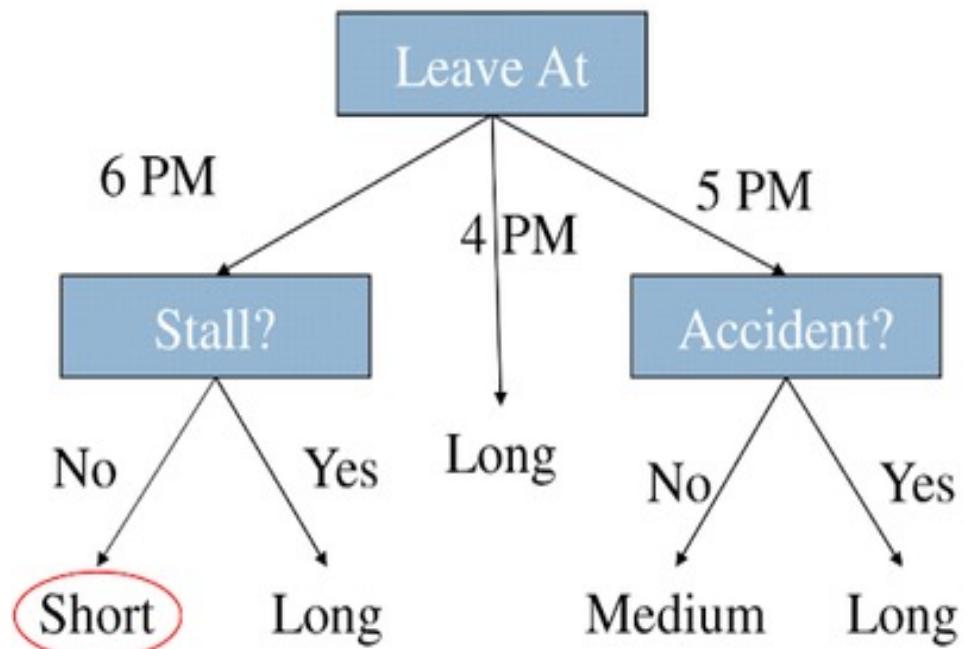
Decision tree algorithms

Model complexity & pruning



Decision Trees

Prediction of Commute Time

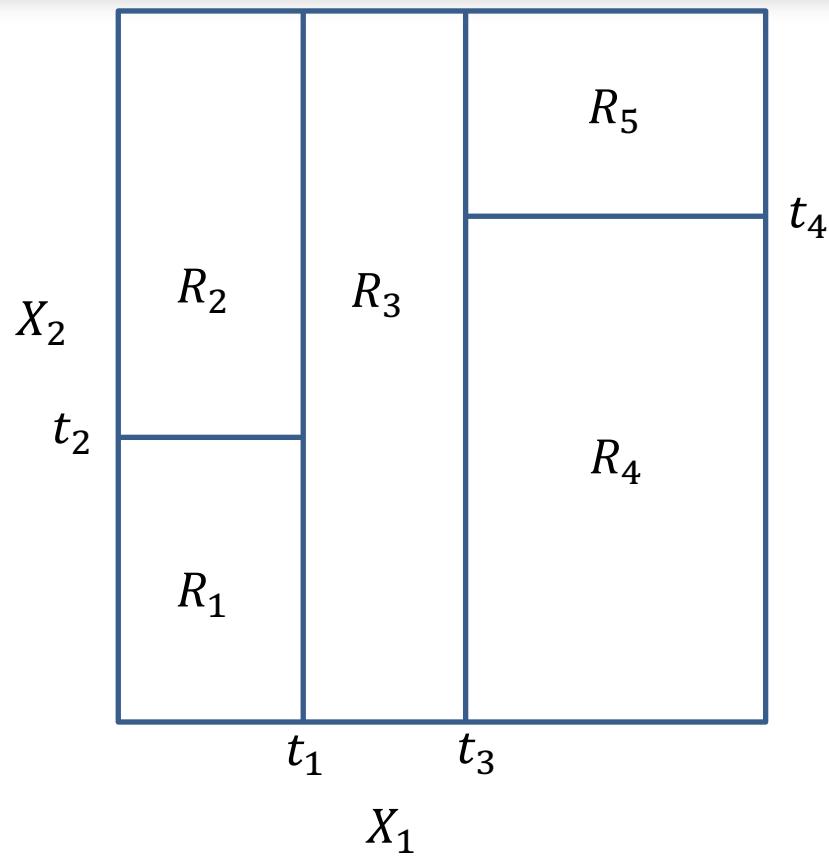
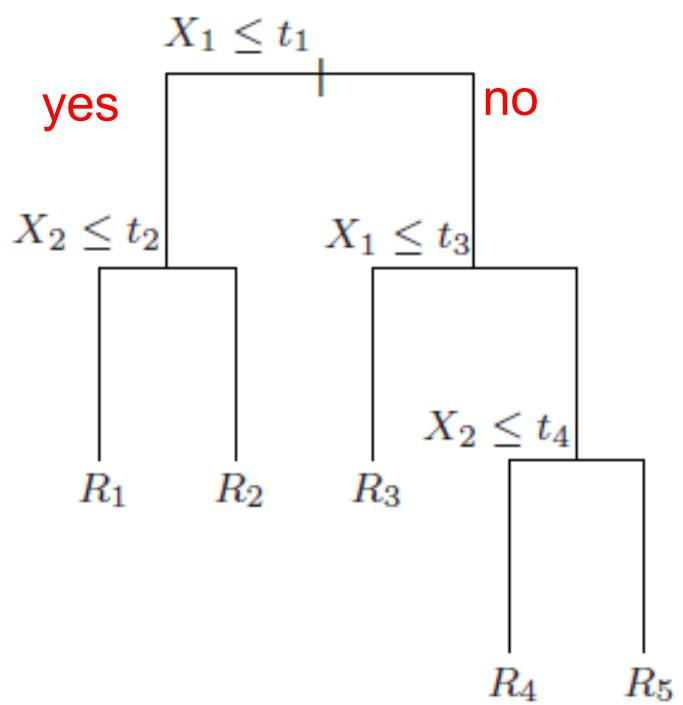


If we leave at 6 PM and there are no cars stalled on the road, what will our commute time be?

Decision Trees

- A decision tree **is a map of the possible outcomes** of a series of related options.
- It **weighs possible actions** against one another
 - Costs, probabilities, and benefits
- Typically starts with a single node
 - Branches into **possible outcomes**.
- **Simple and easy** to interpret!
- **May not be competitive** with the best supervised learning algorithms **in terms of prediction accuracy**.

Partition of Feature Space



[Figure taken from Hastie's
ESL book]

Decision Trees

- After partitioning the feature space, we can fit a simple model in each sub-region ($R_1, R_2 \dots$).
- We can fit a regression model. Such decision trees are called regression trees.
- We can also fit a classification model. Such decision trees are called classification trees.
- Usually, extremely simple models such as majority (classification) or mean (regression) are used.

Process of building a DT

- Let's start with the procedure:
 - We divide the feature space, i.e., the set of possible values for x_1, \dots, x_d into J **distinct and non-overlapping regions**, R_1, \dots, R_J
 - For every instance that falls into region R_j , we make the same prediction, which is simply the **mean (or mode)** of response values for the training observations in R_j .

Formulation of Regression Trees

- The overall goal of regression trees is to find regions R_1, R_2, \dots, R_J that minimize the training error:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

where \hat{y}_{R_j} is the mean of the target values of the training instances in the j^{th} region.

But how do we exactly perform these actions?

Solution

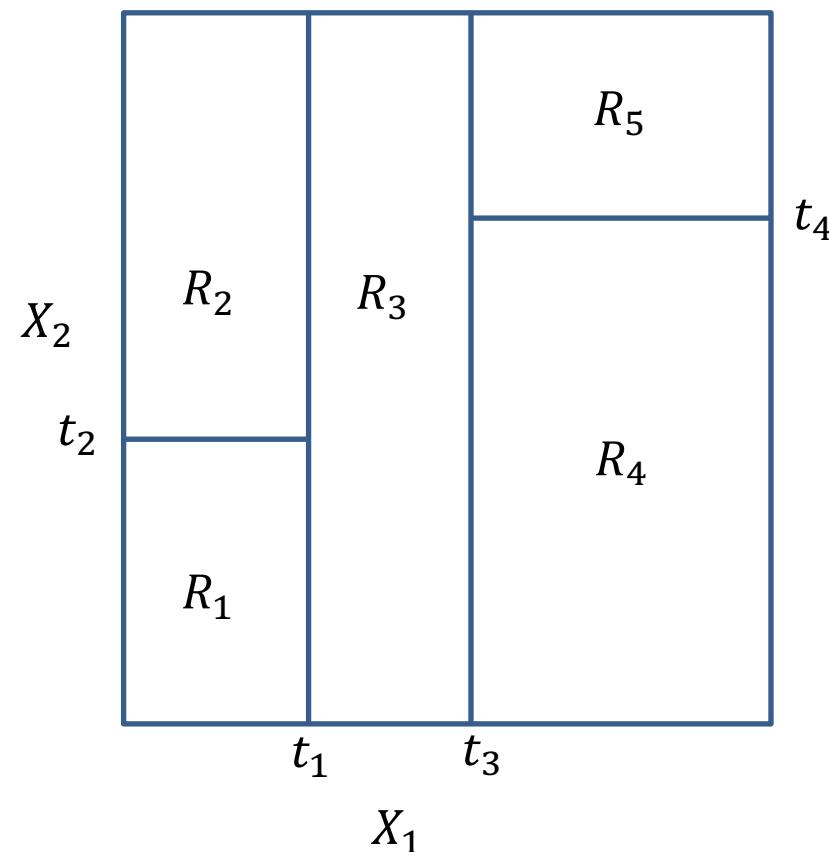
- It is **computationally infeasible** to consider every possible partition of the feature space into J regions.
- For this reason, a **top-down, greedy approach** is used
 - Known as **recursive binary splitting**.
 - Rather than using a **brute-force solution**, we would like to work in a **heuristic** way.

Solution

- How the heuristic method works?
 - Select a feature x_j and a threshold s such that
 - Split the feature space
 - Regions $\{x|x_j \leq s\}$ and $\{x|x_j \geq s\}$
 - leads to the best possible reduction in training error
 - Not going into the joint space of all features
 - Use independent feature form such as x_j with a threshold s .
 - Repeat the process
 - Looking for the best feature and the best threshold
 - Minimize the error in each of the resulting regions.
 - instead of splitting the entire feature space, we only split one of the two previously identified regions.
 - The splitting process continues until a stopping criterion is reached.

Prediction using DT

- We predict the response for a given test instance
 - using the mean (or mode) of the training instances in the region where the test observation falls.



Classification trees...

- In the classification setting
 - we replace the sum of square error by the **classification error rate** as a criterion for making the binary splits.
 - The **classification error rate (E)** is defined as the fraction of the training instances in that region that do not belong to the most common class.

$$E = 1 - \max_k \hat{p}_{jk}$$

where \hat{p}_{jk} represents the proportion (fraction) of training instances in the j^{th} region that are from k^{th} class

$$CoD = \max_k \hat{p}_{jk}$$

- *CoD* (certainty of distribution) and close to 1
 - almost all the training points inside a region are voting for a certain class label.

Classification trees...

- Classification error is being less sensitive for tree-growing
- Alternative solution:
 - **Gini index (G)** is defined as

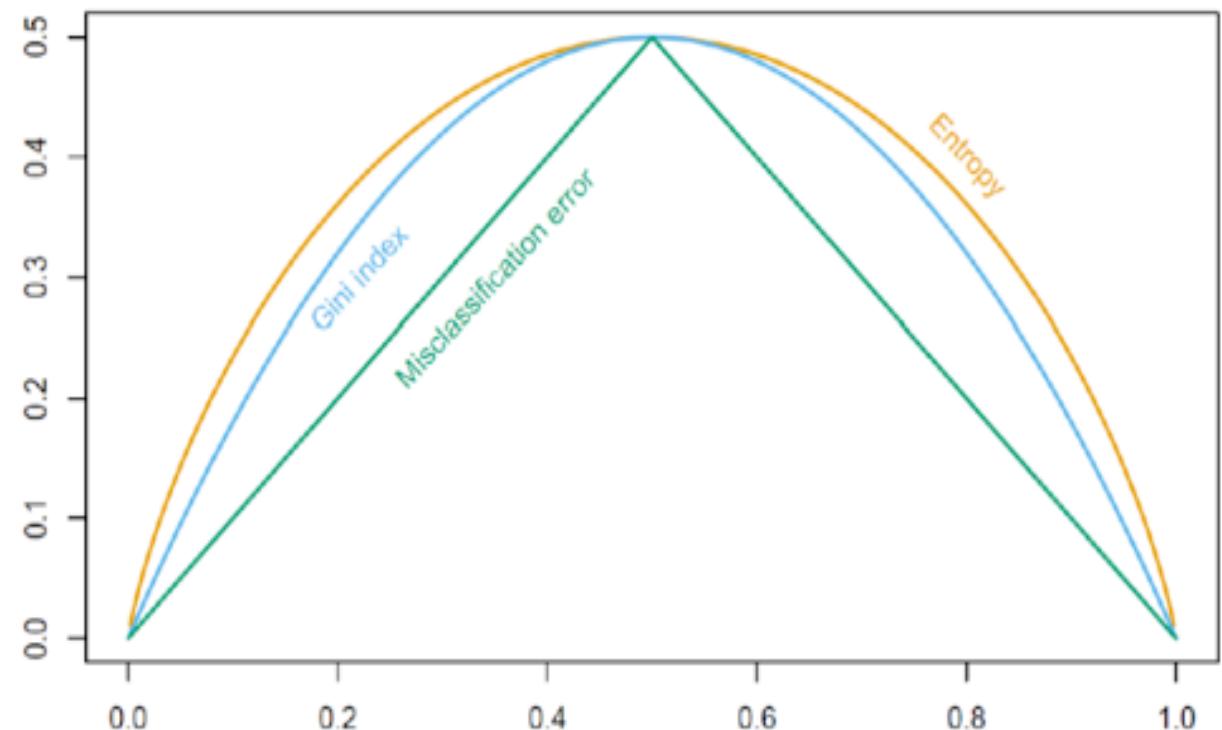
$$G = \sum_{k=1}^K \hat{p}_{jk} (1 - \hat{p}_{jk})$$

- It is a measure of node purity. G becomes small as \hat{p}_{jk} closes to either 0 or 1.
- **Entropy** defined as:

$$D = - \sum_{k=1}^K \hat{p}_{jk} \log \hat{p}_{jk}$$

Classification trees...

- Pattern of Error, Gini Index and Entropy for different probabilities of class distributions:



Decision tree algorithms

- Three of the more popular ones are listed:
 - **ID3** (Iterative Dichotomiser 3)
 - uses **Entropy**
 - **C4.5** (*Successor of ID3*)
 - *slightly more advanced version of ID3 and uses Entropy*
 - **CART** (*Classification and Regression Tree*)
 - uses **Gini impurity**

Decision tree algorithms: ID3 Algorithm...

- Calculate the **entropy** of every feature using the data set S.
- **Split** the set S into subsets **using the feature** for which **entropy** is minimum.
 - So **lesser values of entropy** means it should be **a good choice for selection of the attribute**
- Make a “**decision tree node**” containing that feature.
- **Recurse** on subsets using remaining features.

Decision tree algorithms: ID3 Algorithm...

- If the tree is very deep
 - It partitions the feature space into small regions.
 - Small number of training points in sub-regions.
 - Increases variance and estimation becomes poor
- If the tree is shallow
 - Large regions
 - Small variance but large bias
- Need to find the sweet point
 - Depth of the decision tree
 - Cross-validation as discussed earlier (lecture – week 5)

Model complexity and pruning

- Pruning is a technique that reduces the size of decision trees
 - Removes sections of the tree
 - Little power to classify instances.
- The tree-building process
 - Overfit (creating deep trees)
 - Underfit (creating small number of regions)
- Generally, there are several ways of pruning trees:
 - Pre-pruning (forward pruning)
 - Post-pruning (backward pruning)

Model complexity and pruning: Pre-pruning

- In pre-pruning
 - Decision is made during the building process
 - Stop adding nodes (e.g., by looking at entropy).
- In case of Entropy
 - Check the amount of entropy reduction by selecting different features.
 - Stop splitting when the entropy reduction is not significant.
- Pre-pruning can be problematic
 - Sometimes attributes individually do not contribute much to a decision, but combined, they may have a significant impact.

Model complexity and pruning: Post-pruning

- Post-pruning waits until the full decision tree has been built
 - Then prunes the attributes by subtree replacement.
- Replace an entire subtree with a single region or node
 - It reproduces the smallest error.
- Select a subtree
 - Check – replacing it with a single node or feature incurs a small amount of change in Entropy.
 - If yes, trim the tree. If not, keep that subtree

Decision trees : Advantages/Disadvantages

- Advantages:
 - Decision trees are **very easy to understand**
 - Decision trees are capable of modelling **nonlinear functions**.
 - Decision tree **can handle categorical variable**
- Disadvantages:
 - **Sensitive to small changes** in the data.
 - Adding few data points or change some small values will change the DT
 - **May overfit easily**
 - Deep decision trees increases risk of overfitting and high variance model.
 - **Only axis-aligned splits.**
 - Considers each feature independently
 - No joint probabilities of features
 - **Performance is not competitive**
 - SVM, KNN or Neural network

Impact of distance metrics on KNN performance (Advanced topics)

- KNN classifies new data points according to their closeness
 - Neighbours
 - Distance measures
- Effectiveness of KNN
 - Selection distance metrics
 - Euclidean distance, Manhattan distance, and cosine similarity etc.
- Please use the following link for further explanation.

References:

1. Prasath, V. B., et al. "[Distance and Similarity Measures Effect on the Performance of K-Nearest Neighbor Classifier--A Review.](#)" arXiv preprint arXiv:1708.04321 (2017).

Feature importance of using Decision Trees (DT)

- **Decision Trees** uses feature selection to determine the most important classification features.
- DT operates by recursively segmenting the data into subsets based on the most informative features until a **stopping criterion** is reached.
 - Information gain or the Gini index
- At each node of the tree, the feature with **the highest score** is chosen as the splitting criterion.
 - The significance of each feature can be determined by considering how much it **contributes**.

References:

1. Grabczewski, Krzysztof, and Norbert Jankowski. "[Feature selection with decision tree criterion](#)." Fifth International Conference on Hybrid Intelligent Systems (HIS'05). IEEE, 2005.

Thank You.