



# More Supervised Learning

- Week 10
- Perceptron
- Deep Learning

# This Lecture is about ...


- Motivation to Neural Networks
- Perceptron
- Multi-layer Perceptron (MLP)
- Connections to Deep Learning

# What is different about Neural Networks?

- Linear models **may not be sufficient** when the underlying functions/decision boundaries are extremely **nonlinear**.
- Support vector machines can construct nonlinear functions but use **fixed** feature transformations, which depends on the kernel function.
- Neural Networks allow **the feature transformations to be learnt from data**.




# Historical Motivation

- Our brain has **networks of inter-connected neurons** and has highly-parallel architecture.
- ANN (artificial neural networks) are motivated by **biological neural systems**.
- Two groups of ANN researchers:
  1. Group that uses ANN **to study/model the brain**
  2. Group that uses the brain as the motivation **to design ANN as an effective learning machine**, which might not be the true model of the brain.
- We are following the **second group's approach**.






# Some Examples of ANN






# Perceptron



Motivation &  
Algorithm  
Multi-layer  
Perceptron  
Back-propagation

# Perceptron



- Perceptron is a **simple** neural network used for **binary classification**.
- It has only **one layer with single node**.
  - Given: input vector  $\mathbf{x} = (x_1, x_2, \dots, x_M)$  of  **$M$  dimensions** and weight vector  $\mathbf{w} = (w_0, w_1, \dots, w_M)$
  - The **perceptron produces output**:  $\hat{y} = \text{sign}[v(\mathbf{x}, \mathbf{w})]$  where  $v(\mathbf{x}, \mathbf{w})$  is **the linear combiner**:
$$v(\mathbf{x}, \mathbf{w}) = \sum_{i=1}^M w_i x_i + w_0$$
  - **Better notation**, let  $x_0 = 1$  and  $\mathbf{x} = (x_0, x_1, \dots, x_M)$  then  $v(\mathbf{x}, \mathbf{w}) = \mathbf{w}^T \mathbf{x} = \mathbf{x}^T \mathbf{w}$  and  $\hat{y}(\mathbf{x}, \mathbf{w}) = \text{sign}[\mathbf{w}^T \mathbf{x}]$


# What Perceptron is doing

- Given weight  $\mathbf{w}$ , the perceptron linearly divides input space into two regions:
  - All  $\mathbf{x}$ 's such that  $\hat{y}(\mathbf{x}, \mathbf{w}) = 1$ , or  $v(\mathbf{x}, \mathbf{w}) \geq 0$
  - All  $\mathbf{x}$ 's such that  $\hat{y}(\mathbf{x}, \mathbf{w}) = -1$ , or  $v(\mathbf{x}, \mathbf{w}) < 0$
- This corresponds to the two sides of the hyperplane defined by the equation:


$$v(\mathbf{x}, \mathbf{w}) = \sum_{i=1}^M w_i x_i + w_0 = 0$$

- Note that  $|v(\mathbf{x}, \mathbf{w})|$  is proportional to the distance from  $x$  to the hyperplane.


$$\begin{aligned}\text{dist}(\mathbf{x}, \mathcal{H}(\mathbf{w})) &= \frac{|\sum_{i=1}^M w_i x_i + w_0|}{\sqrt{\sum_{i=1}^M w_i^2}} \\ &= \frac{|v(\mathbf{x}, \mathbf{w})|}{\sqrt{\sum_{i=1}^M w_i^2}}\end{aligned}$$



# What Perceptron is doing



- The sign of  $v(x, \mathbf{w})$  indicates on which side of hyperplane  $\mathcal{H}(\mathbf{w})$  is  $x$ .
- While the magnitude  $|v(x, \mathbf{w})|$  indicates how far away  $x$  is from  $\mathcal{H}(\mathbf{w})$ .




**Linearly separable**

- Two sets  $\mathcal{C}_1$  and  $\mathcal{C}_2$  are called linearly separable if there exists a hyperplane  $\mathcal{H}(\mathbf{w})$  that separates them.

# Training (or Learning) Perceptron

- Find the weight vector  $w$  so that the perceptron correctly classify 2 classes, given sample training data



- Training data  $D = \{(\mathbf{x}_t, y_t)\}, t = 1, \dots, n$  where  $\mathbf{x}_t = (x_{t1}, \dots, x_{tM})$  is the input vector at time  $t$   
 $y_t = \pm 1$  is the desired output

# Learning Perceptron

## Perceptron Learning Algorithm

1. Initialize  $\mathbf{w} = \mathbf{0}$
2. Retrieve next input  $\mathbf{x}_t$  and desired output  $y_t$   
Compute actual output  $\hat{y}_t = \text{sign}[\mathbf{x}_t \cdot \mathbf{w}]$   
Compute output error  $e_t = y_t - \hat{y}_t$   
Update weight, for all  $i$ :  
$$w_i \leftarrow w_i + \Delta w_i \quad \text{with } \Delta w_i = \eta e_t x_{ti}$$
  
where  $0 < \eta \leq 1$  is **learning rate**
3. Repeat from step 2 until convergence

$\mathbf{x}_t = (x_{t1}, \dots, x_{tM})$

**Remark:** no update if  $\hat{y}_t = y_t$ , e.g., the current weight  $\mathbf{w}$  already correctly classify current sample  $\mathbf{x}(t)$

# Learning Perceptron

## Example:

Current weight  $\mathbf{w} = (-1, 2, 1)$

Current training sample  $\mathbf{x} = \left(\frac{1}{2}, 1\right)$ ,  $y = -1$


$$v(\mathbf{x}, \mathbf{w}) = 2 \times \frac{1}{2} + 1 \times 1 - 1 = 1$$

$$\hat{y}(\mathbf{x}, \mathbf{w}) = \text{sign}(1) = 1$$

$$e = y - \hat{y} = -2$$

New

$$\begin{aligned} w_i &= w_i - 2\eta x_i = w_i - 2x_i \quad (\text{let } \eta = 1) \\ w_0 &= -1 - 2x_0 = -3 \\ w_1 &= 2 - 2x_1 = 1 \\ w_2 &= 1 - 2x_2 = -1 \end{aligned}$$




# Learning Perceptron

- Effect of the update rule  
 $\mathbf{w}' \leftarrow \mathbf{w} + \eta e \mathbf{x} = \mathbf{w} + \eta(y - \hat{y})\mathbf{x}$
- Inspect how  $\mathbf{w}'$  classify  $\mathbf{x}$ :  
$$\begin{aligned}\nu' &= \mathbf{w}' \cdot \mathbf{x} = [\mathbf{w} + \eta(y - \hat{y})\mathbf{x}] \cdot \mathbf{x} \\ &= \mathbf{w} \cdot \mathbf{x} + \eta(y - \hat{y})\mathbf{x}^2 \\ &= \nu + \eta(y - \hat{y})\mathbf{x}^2\end{aligned}$$
- Thus, when  $\hat{y} = -1$  but  $y = 1$ , i.e.,  $\nu < 0$  and we want  $\nu > 0$  and indeed  $\nu' > \nu$  (why?)
- When  $\hat{y} = 1$  but  $y = -1$ , i.e.  $\nu \geq 0$  but we want  $\nu < 0$ , and indeed  $\nu' < \nu$

The update rule creates new  $\mathbf{w}'$  that is better than  $\mathbf{w}$  in classifying  $\mathbf{x}$ .  
That's what we want!

# Perceptron Convergence Theorem

- If training instances are drawn from **two linearly separate sets**  $\mathcal{C}_1$  and  $\mathcal{C}_2$ , then the perceptron learning rule **will converge after finite iterations.**




- However, **no guarantee for convergence** if  $\mathcal{C}_1$  and  $\mathcal{C}_2$  are **not linearly separable!**

# Some Linearly Separable Problems


AND

	$x_1$	$x_2$	$y$
a	0	0	0
b	0	1	0
c	1	0	0
d	1	1	1




OR

	$x_1$	$x_2$	$y$
a	0	0	0
b	0	1	1
c	1	0	1
d	1	1	1




# XOR Problem

Perceptron cannot deal with problems such as XOR!



A multilayer Perceptron (MLP) can represent XOR problem



# Learning Perceptron in Python

## sklearn.linear\_model.Perceptron



```
class sklearn.linear_model. Perceptron (penalty=None, alpha=0.0001, fit_intercept=True, n_iter=5, shuffle=True,  
verbose=0, eta0=1.0, n_jobs=1, random_state=0, class_weight=None, warm_start=False) [source]
```

### Methods

<code>decision_function (X)</code>	Predict confidence scores for samples.
<code>densify ()</code>	Convert coefficient matrix to dense array format.
<code>fit (X, y[, coef_init, intercept_init, ...])</code>	Fit linear model with Stochastic Gradient Descent.
<code>fit_transform (X[, y])</code>	Fit to data, then transform it.
<code>get_params ([deep])</code>	Get parameters for this estimator.
<code>partial_fit (X, y[, classes, sample_weight])</code>	Fit linear model with Stochastic Gradient Descent.
<code>predict (X)</code>	Predict class labels for samples in X.
<code>score (X, y[, sample_weight])</code>	Returns the mean accuracy on the given test data and labels.
<code>set_params (*args, **kwargs)</code>	
<code>sparsify ()</code>	Convert coefficient matrix to sparse format.
<code>transform (*args, **kwargs)</code>	DEPRECATED: Support to use estimators as feature selectors will be removed in version 0.19.

# Multi-layer Feed-forward NN

- Perceptron is **quite weak** in what it can represent.
- For **complex, non-linear decision surfaces**, we need multi-layer network.
- Choice of node in multi-layer network
  - Perceptron: **discontinuity**
  - Answer: **sigmoid** function!
- **Sigmoid node:** like a perceptron, but with the sigmoid function  $\sigma(z) = (1 + e^{-z})^{-1}$  instead of the sign function, i.e.,  $y = \sigma(\mathbf{w}^T \mathbf{x})$ .



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\frac{d\sigma(z)}{dz} = \sigma(z)(1 - \sigma(z))$$

# Multi-layer Perceptron


- A feedforward neural network is an ANN wherein **connections between units do not form a cycle**.
- Multi-layer feed-forward NN is **also known as Multi-layer Perceptron (MLP)**.
- The term “MLP” is really a **mismuter**.
  - **Why?** Because the model comprises multiple layers of logistic regression like models (with **continuous nonlinearities**) rather than multiple Perceptrons (with **discontinuous nonlinearities**).
- Although a misnomer, we will continue using **MLP** term.

# Structure of Multi-layer Perceptron

Consider a two-layer network: input layer, hidden layer and output layer

Remarks:

- Output now is a vector
- Two kinds of weights:
  - input → hidden
  - hidden → output
- $w_{ij}^h$ : from  $i^{th}$  input →  $j^{th}$  hidden
- $w_{jk}^o$ : from  $j^{th}$  hidden →  $k^{th}$  output
- Input layer does no computation, only to relay input vector.
- Can have more than one hidden layers.
- Doesn't have to be fully connected.



# MLP Formulation

- Given input  $x_t$  and desired output  $y_t$ ,  $t = 1, \dots, n$ , **find the network weights  $w$**  such that


$$\hat{y}_t \approx y_t, \forall t$$

- Stating above as **an optimization problem**: find  $w$  to minimize the error function

$$E(w) = \frac{1}{2} \sum_{t=1}^n \sum_{k=1}^K (y_{tk} - \hat{y}_{tk})^2$$

↓ sum over all training samples      ↓ sum over all outputs      ↓ error in the  $k^{th}$  output

- We will use **gradient-descent** for minimization.
- $E(w)$  is **not convex**, but a complex function with possibly many local minima.
- We will use an algorithm called **Backpropagation**.



# Detour: Gradient-based Optimization


- To minimize a functional  $f(\mathbf{x})$ , use gradient-descent:
  - Initialize random  $\mathbf{x}_0$
  - Slide down the surface of  $f$  in the direction of steepest decrease:

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta \times \nabla f(\mathbf{x}_t)$$

learning rate

- Similarly, to **maximize**  $f(\mathbf{x})$ , use **gradient-ascent**.

$$\mathbf{x}_{t+1} = \mathbf{x}_t + \eta \times \nabla f(\mathbf{x}_t)$$



# Detour: Stochastic Gradient Descent (SGD)

- Instead of minimizing  $E(\mathbf{w})$ , SGD minimizes the instantaneous approximation of  $E(\mathbf{w})$  using only  $t$ -th instance, i.e.,

$$E_t(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^K (y_{tk} - \hat{y}_{tk})^2$$

- Update rule (where  $t$  denotes the current training sample) is

$$w_i \leftarrow w_i - \eta \frac{\partial E_t(\mathbf{w})}{\partial w_i}$$

- SGD is cheap to perform and guaranteed to reach a local minimum in a stochastic sense.

# Training MLP: Backpropagation

- It is in fact a **stochastic gradient-descent rule!**
- Minimizing instantaneous approximation for current training sample  $(x_t, y_t)$


$$E_t(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^K (\hat{y}_k - y_k)^2$$

- Gradient-descent rule:  $w_{jk}^o \leftarrow w_{jk}^o - \eta \frac{\partial E_t(w)}{\partial w_{jk}^o}$
- Let  $\bar{y}_k$  be the (**unsigmoided**) argument value at output node  $\hat{y}_k$ , i.e.,  $\hat{y}_k = \sigma(\bar{y}_k)$ , we have:

$$\frac{\partial E_t}{\partial w_{jk}^o} = \boxed{\frac{\partial E_t}{\partial \bar{y}_k}} \times \frac{\partial \bar{y}_k}{\partial w_{jk}^o} \quad \text{--- } = z_j \quad \text{since } \bar{y}_k = \sum w_{jk}^o z_j$$

$$\boxed{\frac{\partial E_t}{\partial \bar{y}_k}} = \frac{\partial E_t}{\partial \hat{y}_k} \times \frac{\partial \hat{y}_k}{\partial \bar{y}_k} \quad \left. \begin{array}{l} (1 - \hat{y}_k)\hat{y}_k \\ -(y_k - \hat{y}_k) \end{array} \right\} -\delta_k^o = -(y_k - \hat{y}_k)(1 - \hat{y}_k)\hat{y}_k$$

- This gradient rule implies:  $w_{jk}^o \leftarrow w_{jk}^o + \eta \delta_k^o z_j$



# Training MLP: Backpropagation

- Similarly for input  $\rightarrow$  hidden weights

$$E_t(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^K (\hat{y}_k - y_k)^2 \quad z_j = \sigma(\bar{z}_j) \text{ where } \bar{z}_j = \sum_{i=1}^M x_i w_{ij}^h$$

$$w_{ij}^h \leftarrow w_{ij}^h - \eta \frac{\partial E_t(\mathbf{w})}{\partial w_{ij}^h}$$

$$\frac{\partial E_t(\mathbf{w})}{\partial w_{ij}^h} = \boxed{\frac{\partial E_t}{\partial \bar{z}_j}} \times \frac{\partial \bar{z}_j}{\partial w_{ij}^h}$$

$= x_i$

$$\boxed{\frac{\partial E_t}{\partial \bar{z}_j}} = \frac{\partial E_t}{\partial z_j} \times \frac{\partial z_j}{\partial \bar{z}_j}$$

$= z_j (1 - z_j)$

chain rule

$$\sum_{k=1}^K \boxed{\frac{\partial E_t}{\partial \bar{y}_k}} \times \frac{\partial \bar{y}_k}{\partial z_j}$$


$= w_{jk}^0 \text{ since } \bar{y}_k = \sum w_{jk}^0 z_j$

$$-\delta_k^o$$

$$\frac{\partial E_t}{\partial \bar{z}_j} = -\delta_j^h = -z_j(1 - z_j) \sum_{k=1}^K w_{jk}^o \delta_k^o$$

- Gradient descent update

$w_{ij}^h \leftarrow w_{ij}^h + \eta \delta_j^h x_i$



# Backpropagation (SGD)

## Algorithm

**Input:** training data

Initialize weights // use small random numbers, between -0.5 and 0.5

**Until** stopping criteria is met **do**

**For** each training sample **do**

        // propagate input forward

        Compute values of hidden nodes  $\mathbf{z}$  and output nodes  $\hat{\mathbf{y}}$

        // propagate error backward

        Compute output error term:  $\delta_k^o = \hat{y}_k(1 - \hat{y}_k)(y_k - \hat{y}_k)$

        Compute hidden error term:  $\delta_j^h = z_j(1 - z_j) \sum_{k=1}^K w_{jk}^o \delta_k^o$

        Update weights:

            Hidden  $\rightarrow$  output:  $w_{jk}^o \leftarrow w_{jk}^o + \eta \delta_k^o h_j$

            Input  $\rightarrow$  hidden:  $w_{ij}^h \leftarrow w_{ij}^h + \eta \delta_j^h x_i$

**Output:** trained weights

# Issues with Backpropagation

## Local minima.


- Possible fixes:

- Add a momentum term in the update rule, e.g.,

$$w_{jk}^o \leftarrow w_{jk}^o + \eta \delta_k^o z_j + \alpha \overbrace{\Delta^{t-1} w_{jk}^o}^{\text{momentum constant } > 0}$$

last amount of update of  $w_{jk}^o$


- Can prevent getting stuck in shallow local minimum



- Multiple restarts and choose final network with best performance.

# Issues with Backpropagation

## Overfitting



- The tendency of the network to “memorize” all training samples, leading to poor generalization
- Usually happens with network of too many hidden nodes and overtrained.
- Possible fixes:
  - Use cross validation, e.g., stop training when validation error starts to grow.
  - Weight decaying: minimize also the magnitude of weights, keeping weights small (since the sigmoid function is almost linear near 0, if weights are small, decision surfaces are less non-linear and smoother)
  - Keep small number of hidden nodes!

# Using MLP in Scikit Learn Python

## Methods

<code>fit (X, y)</code>	Fit the model to data matrix X and target y.
<code>get_params ([deep])</code>	Get parameters for this estimator.
<code>predict (X)</code>	Predict using the multi-layer perceptron classifier
<code>predict_log_proba (X)</code>	Return the log of probability estimates.
<code>predict_proba (X)</code>	Probability estimates.
<code>score (X, y[, sample_weight])</code>	Returns the mean accuracy on the given test data and labels.
<code>set_params (**params)</code>	Set the parameters of this estimator.


## `sklearn.neural_network.MLPClassifier`

```
class sklearn.neural_network.MLPClassifier(hidden_layer_sizes=(100,), activation='relu', solver='adam',
alpha=0.0001, batch_size='auto', learning_rate='constant', learning_rate_init=0.001, power_t=0.5, max_iter=200,
shuffle=True, random_state=None, tol=0.0001, verbose=False, warm_start=False, momentum=0.9,
nesterovs_momentum=True, early_stopping=False, validation_fraction=0.1, beta_1=0.9, beta_2=0.999, epsilon=1e-08)
```

[source]

Multi-layer Perceptron classifier.

# Where to go from here?




# Deep Learning

---


- ✓ Deep Learning methods are advanced neural networks.
- ✓ They have been successful in learning many real world tasks  
e.g. handwritten digit recognition, image recognition!
- ✓ Some of the common Deep Learning architectures are:
  - ✓ Convolutional Networks (Due to Le Cun *et al.*)
  - ✓ Autoencoders (Due to Yoshua Bengio *et al.*)
  - ✓ Deep Belief Networks (due to Geoff Hinton *et al.*)
  - ✓ Boltzmann Machines
  - ✓ Restricted Boltzmann Machines
  - ✓ Deep Boltzmann Machines
  - ✓ Deep Neural Networks

# Convolutional Neural Networks


- Also called **CNN** or **ConvNets**.
- Motivation:
  - vision processing in our brain is fast
  - Also (Hubel & Wiesel, 62'):
    - Simple cells detect local features
    - Complex cells **pool** local features
- Translated in technical terms:
  - **Sparse interactions**: sparse weights within a smaller kernel (e.g., 3x3, 5x5) instead of the whole input. **This helps reduce #params.**
  - **Parameter sharing**: a kernel use the same set of weights while applying onto different location (sliding windows).
  - **Translation invariance.**



# Convolutional Neural Networks




# Application of CNN



CIFAR 10 dataset: 50,000 training images, 10,000 test images  
<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

# Peeping into CNN's Brain



# ConvNets won all recent Computer Vision Challenges

- Galaxy



- Plankton



121 classes


- Retina (high-res images)
  - Diabetes recognition from retina image



Read here:

<http://benanne.github.io/2014/04/05/galaxy-zoo.html>

<http://benanne.github.io/2015/03/17/plankton.html>



# What helped Deep Learning?

- Larger models with new training techniques:
  - Dropout, Maxout, Maxnorm, ReLU,...
- Large ImageNet dataset [Fei-Fei et al. 2012]
  - 1.2 million training samples
  - 1000 categories
- Fast graphical processing units (GPU)
  - Capable of 1 trillion operations per second




# Deep Autoencoder

- Simply a neural network that tries to copy its input to its output.

- Input  $x = [x_1, x_2, \dots, x_N]^\top$
- An encoder function  $f$  parameterized by  $\theta$
- A coding representation  $z = [z_1, z_2, \dots, z_K]^\top = f_\theta(x)$
- A decoder function  $g$  parameterized by  $\varphi$
- An output, also called reconstruction  
 $r = [r_1, r_2, \dots, r_N]^\top = g_\varphi(z) = g_\varphi(f_\theta(x))$
- A loss function  $\mathcal{J}$  that computes a scalar  $\mathcal{J}(x, r)$  to measure how good of a reconstruction  $r$  of the given input  $x$ , e.g., mean square error loss:


$$\mathcal{J}(x, r) = \sum_{n=1}^N (x_n - r_n)^2$$

- Learning objective:  $\operatorname{argmin}_{\theta, \varphi} \mathcal{J}(x, r)$



# Deep Autoencoder

- Autoencoders are another way of **feature learning**.
- The solution is **trivial unless there are constraints** (such as **sparsity**) on the number of nodes in the hidden layers.
- Linear autoencoder with  $K \leq N$  acts as **PCA**.
- However, it is **nonlinearity** that makes it powerful.




# Resources

- Courses/Tutorials
  - Hinton Coursera:  
<https://www.coursera.org/course/neuralnets>
  - Hugo Youtube channel:  
<https://www.youtube.com/user/hugolarochelle>
  - Colah's blog (<http://colah.github.io/>): very nice visualizations for easier and better understanding.
  - Karpathy's blog (<http://karpathy.github.io/>): deep learning in web browsers.

# Resources

- Tools

- Tensorflow (Google)
  - Keras
- Theano (Montreal – Bengio et. al.)
  - Lasange, Keras
- Caffe (Berkeley)
- Torch (Facebook)
- Deep Scalable Sparse Tensor Network Engine (DSSTNE) (Amazon)
- MatConvNet (Oxford, MATLAB)



**Thank You.**